
Distributed Training Strategies for Knowledge Base Completion Model(ConvKB)

Vikram Bhatt^{* 1} Kapil Pathak^{* 2}

Abstract

In this report, we present our work on Convolutional Neural Network based knowledge graph embedding model named ConvKB (Nguyen et al., 2017) for knowledge based completion. We trained our model in various distributed training scenarios such as i) Bulk Synchronous Parallel(BSP) Bridging Model ii) Asynchronous Parallel Execution(ASP) iii) DOWNPOUR SGD(Dean et al., 2012b)- an asynchronous stochastic gradient procedure supporting large number of replicas in Distributed Tensorflow framework (Abadi et al., 2015). We initially performed experiments on benchmark dataset FB15K-237 and explored scalability issues in these distributed settings for link prediction tasks. We pre-processed entire Freebase(RDF triples) dump using Hadoop, for the purpose of demonstration of scalability, we constructed our own dataset(FB-1GB). Experiments shows that DOWNPOUR gives best trade-off between training time and convergence stability. Although, there doesn't exist explicit API's in Tensorflow at the moment, for Stale Synchronous Parallel(SSP) Bridging Model, we propose (one of many possible) simple method for implementation in the framework. Our vanilla version of distributed ConvKB model achieves competitive results on link prediction performance on benchmark dataset FB15k-237.

1. Introduction

In this report, the first section talks about what Freebase dataset, its ontology, general schema. Given that it's a

^{*}Equal contribution ¹Department of Computational and Data Sciences, Indian Institute of Science, Bangalore ²Department of Computer Science and Automation, Indian Institute of Science, Bangalore. Correspondence to: Vikram Bhatt <vikrambhatt@iisc.ac.in>, Kapil Pathak <kapilpathak@iisc.ac.in>.

huge dataset, this section talks about the preprocessing techniques we have implemented. The second section talks about ConvKB model. The third section talks about various distributed learning schemes we have tried and their analysis. The fourth section talks about results and analysis. The last section is dedicated for the further ideas developed during the project and are yet to be tried.

2. Literature Review

Knowledge graphs are graph structured bases, where facts are represented as triplets consisting subject entity (head), relation and object entity (tail) given as (h,r,t). However, knowledge bases tend to suffer from incompleteness and identifying missing links is referred to as link prediction. Considering the fact, these bases contain billions of facts, link prediction should scale well with respect to parameters of our model and computational costs. Various embedding methods have been proposed previously (Bordes et al., 2013) (Trouillon et al., 2016) (Nguyen et al., 2016) using translational based characteristics on low dimensional embeddings of relations and entities. Although these models are simple and scalable, they learn less expressive features. A simple way to learn more expressive features capturing inter-relationships between entities and relations, is to use deep learning models. Convolutional neural networks tend to capture non-linear relations between the embeddings with few number of parameters. (Dettmers et al., 2018) proposed ConvE model, a multilayer convolutional network which is parameter efficient and achieved state of the art results on various benchmark datasets. (Nguyen et al., 2017) in their article state that,

ConvE doesn't capture global relationships among same dimensional entries of an embedding triple, so ConvE ignores the transitional characteristic in transition-based models, which is one of the most useful intuitions for the task. ConvKB models relationships among same dimensional entries of the embeddings.

3. Model-ConvKB

We denote a valid factual triple in the form of $(head, relation, tail)$ as (h, r, t) and the corresponding k -dimensional embeddings as (v_h, v_r, v_t) . We stack these embedding horizontally in a matrix form $A = [v_h, v_r, v_t] \in R^{k \times 3}$. We use τ number of filters $\omega \in R^{1 \times 3}$, which slides on each row of matrix of A , as illustrated in the Fig.1 and generate k feature maps $v = [v_1, v_2, \dots, v_k] \in R^k$ where $v_i = g(\omega * A_{i,:} + b)$, $b \in R$ is bias and g is non-linear function (such as ReLU). These τ feature maps concatenated vertically and formed single vector of shape $R^{\tau k \times 1}$, then computed dot product with vector $w \in R^{k \times 1}$ of same size to give score for triple (h, r, t) .

$$f(h, r, t) = \text{concat}(g([v_h, v_r, v_t] * \Omega)).w$$

where f is score function.

We use loss function L with L_2 regularization on the weight vector w

$$L = \sum_{(h,r,t) \in \{G \cup G'\}} \log(1 + \exp(l_{(h,r,t)} \cdot f(h, r, t))) + \frac{\lambda}{2} \|w\|^2$$

where,

$$l_{(h,r,t)} = \begin{cases} 1 & \text{for } (h, r, t) \in G, \\ -1 & \text{for } (h, r, t) \in G' \end{cases}$$

G' is a collection of invalid tuples generated by corrupting valid tuples in G . We use the Bernoulli Trick to generate head or tail for sampling invalid triples. (Wang et al., 2014). As we need to scale up this model for larger dataset, we add 1 more layer of CNN layer on the top of initial layer.

4. Experiments

4.1. Bulk Synchronous Parallel(BSP)

In Bulk Synchronous Parallel (BSP), the model alternate between computation phase and communication phase or also called synchronized barrier. Each worker accumulates the gradients for all parameters until the next superstep and the effect of the computations of one worker node is invisible to other worker nodes until the next superstep occurs. Because of the cleaner separation of computation and communication phase, BSP scheme often enjoys correctness of parallel machine learning algorithms.

In (Abadi et al., 2016) Tensorflow API, *SyncReplicasOptimizer* performs similar job to BSP. Here all accumulators, each for individual parameter performs the job of gradients average. After each global step, the accumulator checks the staleness of all gradients and drops the gradients before those. A drawback of this scheme is that the throughput of individual worker is severely affected as each worker needs to wait until the next superstep occurs. Fig.2 and Fig.3 shows




Figure 1. Graphical illustration of ConvKB model (Nguyen et al., 2017) with embedding size $k = 4$, the number of filters $\tau = 3$.

Table 1. Training times for various number of workers with one (fixed) parameter server.

NO. OF WORKERS	WALLCLOCK TIME
2	7HR 12M
3	6HR 4M
4	3HR 12M
5	3HR

the learning curve for BSP training with 2 and 5 workers respectively. Table 1 shows training times on FB15k-237 dataset.

4.2. Asynchronous SGD

Unlike BSP, in asynchronous SGD, any worker never waits for any superstep or other worker to catch synchronize the computations. Each worker computes gradients and communicates with parameter server separately throughout the course of each iteration. This scheme achieves near-ideal P-fold throughput in each iteration. But as there is no synchronization among the co-workers, the staleness of the gradients causes slower convergence rate per iteration than BSP scheme. If the staleness is not bounded, it may even cause incorrect gradient calculations and wrong results. This is a drawback of asynchronous SGD.



Figure 2. Learning curve for synchronous SGD with 2 workers and 1 parameter server.




Figure 4. Learning curve for asynchronous SGD with 2 workers and 1 parameter server.



Figure 3. Learning curve for synchronous SGD with 5 workers and 1 parameter server



Figure 5. Learning curve for asynchronous SGD with 5 workers and 1 parameter server.

4.3. Stale Synchronous Parallel (SSP) Bridging model

In case of asynchronous SGD, as unbounded staleness may cause incorrect results, Stale Synchronous Parallel (SSP) allows us to overcome this issue by bounding the staleness for particular s . For given staleness s , SSP allows a faster worker go ahead in terms of iterations than its other co-workers until it is s iterations apart. The local model parameters become stale as each worker receives partial updates. The advantage of SSP model is that it behaves like Asynchronous model most of the time. It can also stop workers as needed to ensure correct update of parameters. SSP bridging model finds advantages of both BSP and asynchronous SGD as it increases throughput per iteration compared to BSP and avoids unbounded staleness, enhancing convergence rate. (Xing et al., 2016)

4.4. Downpour SGD

Another variant of asynchronous stochastic gradient descent that can be applied on large dataset is Downpour SGD. Here too, the model replicas run independently along with independent parameter server shards. In Downpour SGD, the communication overhead can be reduced by limiting parameter requests from each replica to parameter server after n_{push} steps and n_{fetch} steps vice versa. In Downpour SGD, a separate Adagrad optimizer is applied on the local copy of parameters on each worker. Here each model replica is calculating gradients on the set of parameters which are stale depending upon the n_{fetch} and n_{push} steps. This introduces stochasticity into the model. (Dean et al., 2012a) suggest that Adagrad optimizer along with "warmstarting" strategy allows to keep more number of replicas that can work productively and serve the purpose of distributed train-



Figure 6. Learning curve for SSP-SGD with 4 workers and 1 parameter server.




Figure 7. Learning curve for SSP-SGD with 6 workers and 1 parameter server.

ing. This induces more stability into the training of deep networks. Published ConvKB model uses single layer of CNN. But if we need to scale up the knowledge graph embeddings further, we can add more number of CNN layers to the model which induces more nonlinearity into the model. Hence, apart from Stale Synchronous Parallel (SSP), Downpour SGD is an effective distributed learning strategy, especially for ConvKB model.

5. Freebase

The Freebase knowledge graph is an important semantic web and linked data technology. It was acquired by Google in 2010 and shut down in 2016. This entire data dump is of size 33 GB in compressed(gZip) format. This data is useful for various research tasks such as information retrieval, knowledge based question answering etc.

5.1. Freebase Schema

In Freebase, each entity is represented by unique machineId identifier. Each machineId starts with "/m/" followed by alphanumeric characters. Each relation is represented by domain followed by type followed by property and each of them separated by "/". For example, for machineId "/m/xyz789" the place of birth given by machineId "/m/cincinnati123" can be incorporated into triplet as follows

/m/xyz789, /people/person/place_of_birth,

/m/cincinnati123

Here machineId /m/xyz789 is of domain people, of type person and a machineId /m/cincinnati123 is related to /m/xyz789 with property place_of_birth. Such kind of triplets are called facts. Within each type of domain, there are further finer-grained details given by /property. The following examples states the fact /m/syz789 entity is part of topic people and has type property person.

/m/xyz789, /type/object/type, /people/person

5.2. Preprocessing of Freebase Dataset

This dataset is in the form of N-Triples RDF format, each of the entities of the format "/m/xyz789" has been converted to "/m.xyz789" while each of the relations of the format "/people/person" has been converted to "/people.person". These changes have reverted to original form by replacing "." by "/". Apart from this change, each object in the data set is encoded with full URL path such as

http://rdf.freebase.com/ns/m.xyz789

In the preprocessing step, these URLs are removed to reduce the size of the dataset as these URLs are not required to incorporate into the model. Apart from machineIds, the triplets consisting of *mediators* starting with "/g" are also present in the data dumps. These triplets consisting of *mediators* only convey properties of the head entity. We reduced the entire

Algorithm 1 DOWNPOURSGDCLIENT(Dean et al., 2012a)(α, n_fetch, n_push)

procedure

StartAsynchronouslyFetchingParameters(parameters):
parameters=GetParametersFromParamServer()

procedure

StartAsynchronouslyPushingGradients(accruedgradients):
SendGradientStopParamServer(accruedgradients)
accruedgradients = 0

global parameters, accruedgradients

step = 0

accruedgradients = 0

while true do

if (step mod n_fetch) == 0 **then**

StartAsynchronouslyFetchingParameters(parameters)

end if

data =GetNextMinibatch()

gradient = ComputeGradient(parameters, data)

accruedgradients = accruedgradients + gradient

parameters = parameters - $\alpha \times$ gradient

end while

if (step mod n_push) == 0 **then**

StartAsynchronouslyPushingParameters(parameters)

end if

dataset of Freebase to only facts in the form of

/m/070xg/sports/sports_team/colors/m/01g5v

with Hadoop map-reduce job.

5.3. Characteristics of Freebase Dataset

Fig.12 shows the distribution of number of triples topics obtained in the reduced dataset obtained after mapreduce job. We observe the distribution is very skewed, if we want to sample using these proportions one might not sample any triplets from the tail at all. We want to create a new train, test and valid sets from the dataset we obtained from mapreduce job, one can take two ways.

1. Deterministic Splitting

- Since our dataset is huge(17GB) shuffle, sort and split in memory is costly, we decided to go for probabilistic splitting.

2. Probabilistic Splitting

- We stream through the entire dataset and generate a random number uniformly between 0 and

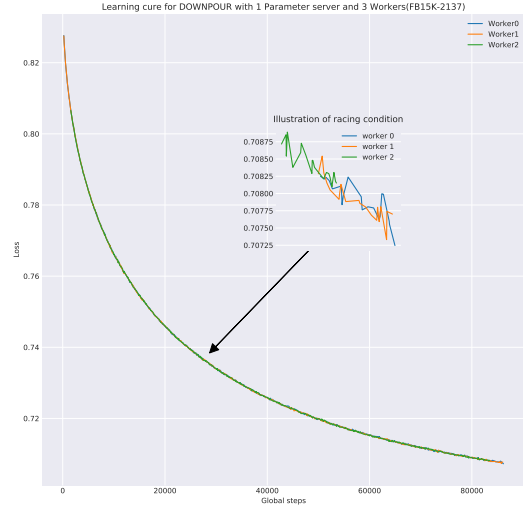


Figure 8. Learning curve for Downpour SGD with 3 workers and 1 parameter server.

Table 2. Comparison for various Benchmarks with and without distributed frameworks on FB15k-237

VARIANTS	MR	MRR	HITS@10
CONVE	246	0.316	49.1
CONVKB	257	0.396	51.7
BSP CONVKB	278.16	0.332	46.7
ASP CONVKB	344.46	0.288	43.4
SSP CONVKB	342.43	0.307	44.2
DOWNPOUR CONVKB	312	0.301	42.6

1.If the generated random number is greater than specified threshold (threshold specifies the percentage of triples in original dataset we want to retain) we add that to our new dataset. This kind of splitting is cheap and takes care of shuffling by the randomness and we never run out of memory.

6. Test Protocol

We used mean rank(MR), mean reciprocal rank(MRR) and Hits@10 as evaluation metrics for our model. Table.2 shows our results for various model we trained. We are using FB15k dataset for these evaluations.

Acknowledgements

Do not include acknowledgements in the initial version of the paper submitted for blind review.




Figure 9. Learning curve for Downpour SGD with 4 workers and 1 parameter server.




Figure 10. Learning curve for Downpour SGD with 5 workers and 1 parameter server.

If a paper is accepted, the final camera-ready version can (and probably should) include acknowledgements. In this case, please place such acknowledgements in an unnumbered section at the end of the paper. Typically, this will include thanks to reviewers who gave useful comments, to colleagues who contributed to the ideas, and to funding agencies and corporate sponsors that provided financial support.

References

- Abadi, Martín, Agarwal, Ashish, Barham, Paul, Brevdo, Eugene, Chen, Zhifeng, Citro, Craig, Corrado, Greg S., Davis, Andy, Dean, Jeffrey, Devin, Matthieu, Ghemawat, Sanjay, Goodfellow, Ian J., Harp, Andrew, Irving, Geoffrey, Isard, Michael, Jia, Yangqing, Jozefowicz, Rafal, Kaiser, Lukasz, Kudlur, Manjunath, Levenberg, Josh, Mané, Dan, Monga, Rajat, Moore, Sherry, Murray, Derek, Olah, Chris, Schuster, Mike, Shlens, Jonathon, Steiner, Benoit, Sutskever, Ilya, Talwar, Kunal, Tucker, Paul, Vanhoucke, Vincent, Vasudevan, Vijay, Viégas, Fernanda B., Vinyals, Oriol, Warden, Pete, Wattenberg, Martin, Wicke, Martin, Yu, Yuan, and Zheng, Xiaoqiang. TensorFlow: Large-scale machine learning on heterogeneous systems, 2015. URL <http://tensorflow.org/>. Software available from tensorflow.org.
- Abadi, Martín, Agarwal, Ashish, Barham, Paul, Brevdo, Eugene, Chen, Zhifeng, Citro, Craig, Corrado, Gregory S., Davis, Andy, Dean, Jeffrey, Devin, Matthieu, Ghemawat, Sanjay, Goodfellow, Ian J., Harp, Andrew, Irving, Geoffrey, Isard, Michael, Jia, Yangqing, Józefowicz, Rafal, Kaiser, Lukasz, Kudlur, Manjunath, Levenberg, Josh, Mané, Dan, Monga, Rajat, Moore, Sherry, Murray, Derek Gordon, Olah, Chris, Schuster, Mike, Shlens, Jonathon, Steiner, Benoit, Sutskever, Ilya, Talwar, Kunal, Tucker, Paul A., Vanhoucke, Vincent, Vasudevan, Vijay, Viégas, Fernanda B., Vinyals, Oriol, Warden, Pete, Wattenberg, Martin, Wicke, Martin, Yu, Yuan, and Zheng, Xiaoqiang. Tensorflow: Large-scale machine learning on heterogeneous distributed systems. *CoRR*, abs/1603.04467, 2016. URL <http://arxiv.org/abs/1603.04467>.
- Bordes, Antoine, Usunier, Nicolas, Garcia-Duran, Alberto, Weston, Jason, and Yakhnenko, Oksana. Translating embeddings for modeling multi-relational data. In Burges, C. J. C., Bottou, L., Welling, M., Ghahramani, Z., and Weinberger, K. Q. (eds.), *Advances in Neural Information Processing Systems 26*, pp. 2787–2795. Curran Associates, Inc., 2013. URL <http://papers.nips.cc/paper/5071-translating-embeddings-for-modeling-multi-rel.pdf>.
- Dean, Jeffrey, Corrado, Greg, Monga, Rajat, Chen, Kai, Devin, Matthieu, Mao, Mark, aurelio Ranzato, Marc, Senior, Andrew, Tucker, Paul, Yang, Ke, Le, Quoc V., and Ng, Andrew Y. Large scale distributed deep networks. In




Figure 11. Normalized Wall Clock Time for worker0 with different number of worker nodes




Figure 12. Frequency of topics processed from the entire Free-base dump shows very skew distribution, dominated by few topics weighing heavily compared to the rest.

Pereira, F., Burges, C. J. C., Bottou, L., and Weinberger, K. Q. (eds.), *Advances in Neural Information Processing Systems 25*, pp. 1223–1231. Curran Associates, Inc., 2012a. URL <http://papers.nips.cc/paper/4687-large-scale-distributed-deep-networks.pdf>.

Dean, Jeffrey, Corrado, Greg, Monga, Rajat, Chen, Kai, Devin, Matthieu, Mao, Mark, Senior, Andrew, Tucker, Paul, Yang, Ke, Le, Quoc V, et al. Large scale distributed deep networks. In *Advances in neural information processing systems*, pp. 1223–1231, 2012b.

Dettmers, Tim, Pasquale, Minervini, Pontus, Stenetorp, and Riedel, Sebastian. Convolutional 2d knowledge graph embeddings. In *Proceedings of the 32th AAAI Conference on Artificial Intelligence*, pp. 1811–1818, February 2018. URL <https://arxiv.org/abs/1707.01476>.

Nguyen, Dai Quoc, Nguyen, Tu Dinh, Nguyen, Dat Quoc, and Phung, Dinh. A novel embedding model for knowledge base completion based on convolutional neural network. *arXiv preprint arXiv:1712.02121*, 2017.

Nguyen, Dat Quoc, Sirts, Kairit, Qu, Lizhen, and Johnson, Mark. Stranse: a novel embedding model of entities and relationships in knowledge bases. *CoRR*, abs/1606.08140, 2016. URL <http://arxiv.org/abs/1606.08140>.

Trouillon, Théo, Welbl, Johannes, Riedel, Sebastian, Gaussier, Éric, and Bouchard, Guillaume. Complex embeddings for simple link prediction. In *International Conference on Machine Learning*, pp. 2071–2080, 2016.

Wang, Zhen, Zhang, Jianwen, Feng, Jianlin, and Chen,

Zheng. Knowledge graph embedding by translating on hyperplanes. In *AAAI*, volume 14, pp. 1112–1119, 2014.

Xing, Eric P, Ho, Qirong, Xie, Pengtao, and Wei, Dai. Strategies and principles of distributed machine learning on big data. *Engineering*, 2(2):179–195, 2016.

A. Acknowledgements

1. We would like to thank Dai Quoc Nguyen for providing baseline model and helper functions. (<https://github.com/daiquocnguyen/ConvKB>)
2. We would like to appreciate Tommy Mulc (<https://github.com/tmulc18/Distributed-TensorFlow-Guide>) for their awesome tutorials on distributed tensorflow frameworks.